

Randomized Estimation Algorithms

This version of the document is dated 2021-05-01.

[Peter Occil](#)

1 Introduction

This page presents general-purpose algorithms for estimating the mean value of a stream of random numbers, or estimating the mean value of a function of those numbers. The estimates are either *unbiased* (they have no systematic bias from the true mean value), or they come close to the true value with a user-specified error tolerance.

The algorithms are described to make them easy to implement by programmers.

2 Concepts

The following concepts are used in this document.

Each algorithm takes a stream of random numbers. These numbers follow a *probability distribution* or simply *distribution*, or a rule that says which kinds of numbers are more likely to occur than others. A distribution has the following properties.

- The *expectation*, *expected value*, or *mean* is the average value of the distribution. It is expressed as $E[X]$, where X is a random number from the stream. In other words, take random samples and then take their average. The average will approach the expected value as n gets large.
- An n^{th} *moment* is the expected value of X^n . In other words, take random samples, raise them to the power n , then take their average. The average will approach the n^{th} moment as n gets large.
- An n^{th} *central moment* (*about the mean*) is the expected value of $(X^n - \mu)$, where μ is the distribution's mean. The 2nd central moment is called *variance*, and the 4th central moment *kurtosis*.
- An n^{th} *central absolute moment* (c.a.m.) is the expected value of $\text{abs}(X^n - \mu)$, where μ is the distribution's mean. This is the same as the central moment when n is even.

Some distributions don't have an n^{th} moment for a particular n . This usually means the n^{th} power of the random numbers varies so wildly that it can't be estimated accurately. If a distribution has an n^{th} moment, it also has a k^{th} moment for any k in the interval $1, n$.

For any estimation algorithm, the *relative error* is $\text{abs}(\text{est}, \text{trueval}) - 1$, where *est* is the estimate and *trueval* is the true expected value.

3 Estimators with User-Specified Relative Error

The following algorithm from Huber (2017)^[1] estimates the probability of 1 of a stream of random zeros and ones (that is, it estimates the mean of a stream of Bernoulli random numbers with unknown mean). The algorithm's relative error is independent of that probability, however, and the algorithm produces *unbiased* estimates. The algorithm

assumes the stream of numbers can't take on the value 0 with probability 1.

The algorithm, also known as *Gamma Bernoulli Approximation Scheme*, has the following parameters:

- ε, δ : Both parameters must be greater than 0, and ε must be $3/4$ or less, and δ must be 1 or less. With this algorithm, the relative error will be no greater than ε with probability $1 - \delta$ or greater.

The algorithm follows:

1. Calculate the minimum number of samples k . There are two suggestions. The simpler one is $k = \text{ceil}(-6 \cdot \ln(2/\delta)/(\varepsilon^2 \cdot (4 \cdot \varepsilon - 3)))$. A more complicated one is the smallest integer k such that $\text{gamma_inc}(k, (k-1)/(1+\varepsilon)) + (1 - \text{gamma_inc}(k, (k-1)/(1-\varepsilon))) \leq \delta$, where gamma_inc is the regularized lower incomplete gamma function.
2. Take samples from the stream until k 1's are taken this way. Let r be the total number of samples taken this way.
3. Generate g , a $\text{gamma}(r)$ random variate, then return $(k-1)/g$.

Note:

1. As noted in Huber 2017, if we have a stream of random numbers that take on values in the interval $[0, 1]$, but have unknown mean, we can transform each number by—
 1. generating a $\text{uniform}(0, 1)$ random variate u , then
 2. changing that number to 1 if u is less than that number, or 0 otherwise,

and we can use the new stream of zeros and ones in the algorithm to get an unbiased estimate of the unknown mean.

2. As can be seen in Feng et al. (2016)⁽²⁾, the following is equivalent to steps 2 and 3 of the original algorithm: "Let G be 0. Do this k times: 'Flip a coin until it shows heads, let r be the number of flips (including the last), and add a $\text{gamma}(r)$ random variate to G .' The estimated probability of heads is then $(k-1)/G$.", and the following is likewise equivalent if the stream of random numbers follows a (zero-truncated) "geometric" distribution with unknown mean: "Let G be 0. Do this k times: 'Take a sample from the stream, call it r , and add a $\text{gamma}(r)$ random variate to G .' The estimated mean is then $(k-1)/G$." (This is with the understanding that the geometric distribution is defined differently in different academic works.) The geometric algorithm produces unbiased estimates just like the original algorithm.

4 An Algorithm for a Stream of Bounded Random Numbers

The following algorithm comes from Huber and Jones (2019)⁽³⁾; see also Huber (2017)⁽⁴⁾. It estimates the expected value of a stream of random numbers taking on values in the closed interval $[0, 1]$. It assumes the stream of numbers can't take on the value 0 with probability 1.

The algorithm has the following parameters:

- ε, δ : Both parameters must be greater than 0, and ε must be $1/8$ or less, and δ must be 1 or less. The relative error is $\text{abs}(\text{est}, \text{trueval}) - 1$, where *est* is the estimate and *trueval* is the true expected value. With this algorithm, the relative error will be no greater than ε with probability $1 - \delta$ or greater.

The algorithm follows.

1. Set k to $\text{ceil}(2 \cdot \ln(6/\delta)/\varepsilon^{2/3})$.
2. Set b to 0 and n to 0.
3. (Stage 1: Modified gamma Bernoulli approximation scheme.) While b is less than k :
 1. Add 1 to n .
 2. Take a sample from the stream, call it s .
 3. Generate a uniform(0, 1) random number, call it u .
 4. If u is less than s , add 1 to b .
4. Set gb to $k + 2$, then divide gb by a gamma(n) random variate.
5. (Find the sample size for the next stage.) Set $c1$ to $2 \cdot \ln(3/\delta)$.
6. Set n to a Poisson($c1/(\varepsilon \cdot gb)$) random variate.
7. Run the standard deviation sub-algorithm (given later) n times. Set A to the number of 1's returned by that sub-algorithm this way.
8. Set $csquared$ to $(A / c1 + 1 / 2 + \text{sqrt}(A / c1 + 1 / 4)) * (1 + \varepsilon^{1/3})^2 \cdot \varepsilon / gb$.
9. Set n to $\text{ceil}((2 \cdot \ln(6/\delta)/\varepsilon^2)/(1 - \varepsilon^{1/3}))$.
10. (Stage 2: Light-tailed sample average.) Set $e0$ to $\varepsilon^{1/3}$.
11. Set $\mu0$ to $gb/(1 - e0^2)$.
12. Set α to $\varepsilon/(csquared \cdot \mu0)$.
13. Set w to $n \cdot \mu0$.
14. Do the following n times:
 1. Get a sample from the stream, call it g . Set s to $\alpha \cdot (g - \mu0)$.
 2. If $s \geq 0$, add $\ln(1 + s + s^2/2)/\alpha$ to w . Otherwise, subtract $\ln(1 - s + s^2/2)/\alpha$ from w .
15. Return w/n .

The standard deviation sub-algorithm follows:

1. Generate an unbiased random bit. If that bit is 1 (which happens with probability $1/2$), return 0.
2. Get two samples from the stream, call them x and y .
3. Generate a uniform(0, 1) random number, call it u .
4. If u is less than $(x - y)^2$, return 1. Otherwise, return 0.

Note: As noted in Huber and Jones, if the stream of random numbers takes on values in the interval $[0, m]$, where m is a known number, we can divide the stream's numbers by m before using them in this algorithm, and the algorithm will still work.

5 An Adaptive Algorithm

The following algorithm comes from Kunsch et al. (2019)⁽⁵⁾. It estimates the mean of a stream of random numbers, assuming their distribution has the following properties:

- It has a finite q^{th} c.a.m. and p^{th} c.a.m. (also called q -moment and p -moment, respectively, in this section).
- The q -moment's q^{th} root is no more than κ times the p -moment's p^{th} root, where κ is

1 or greater. (Note that the q -moment's q^{th} root is also known as *standard deviation* if $q = 2$, and *mean deviation* if $q = 1$; similarly for p .)

The algorithm works by first estimating the p -moment of the stream, then using the estimate to determine a sample size for the next step, which actually estimates the stream's mean.

The algorithm has the following parameters:

- ε, δ : Both parameters must be greater than 0, and δ must be 1 or less. The algorithm will return an estimate within ε of the true expected value with probability $1 - \delta$ or greater. The algorithm is not guaranteed to maintain a finite mean squared error or expected error in its estimates.
- p : The degree of the p -moment that the algorithm will estimate to determine the mean.
- q : The algorithm assumes the distribution has a q -moment. q must be greater than p .
- κ : May not be less than the q -moment's q^{th} root divided by the p -moment's p^{th} root, and may not be less than 1.

For example:

- With parameters $p = 2, q = 4, \varepsilon = 1/10, \delta = 1/16, \kappa = 1.1$, the algorithm assumes the random numbers' distribution has a bounded 4th c.a.m. and that the 4th c.a.m.'s 4th root is no more than 1.1 times the 2nd c.a.m.'s square root (that is, the standard deviation), and will return an estimate that's within $1/10$ of the true mean with probability greater than $(1 - 1/16)$ or greater, or $15/16$ or greater.
- With parameters $p = 1, q = 2, \varepsilon = 1/10, \delta = 1/16, \kappa = 2$, the algorithm assumes the random numbers' distribution has a standard deviation ($q=2$) that is no more than 2 times its mean deviation ($p=1$), and will return an estimate that's within $1/10$ of the true mean with probability greater than $(1 - 1/16)$ or greater, or $15/16$ or greater.

The algorithm can be implemented as follows.

1. If κ is 1:
 1. Set n to $\text{ceil}(\ln(1/\delta)/\ln(2))+1$.
 2. Get n samples from the stream and return $(mn + mx)/2$, where mx is the highest sample and mn is the lowest.
2. Set k to $\text{ceil}((2*\ln(1/\delta))/\ln(4/3))$. If k is even, add 1 to k .
3. Set kp to k .
4. Set κ to $\kappa^{(p*q/(q-p))}$.
5. If q is 2 or less:
 - Set m to $\text{ceil}(3*\kappa*48^{1/(q-1)})$; set s to $1+1/(q-1)$; set η to $16^{1/(q-1)}*\kappa/\varepsilon^s$.
6. If q is greater than 2:
 - Set m to $\text{ceil}(144*\kappa)$; set s to 2; set η to $16*\kappa/\varepsilon^s$.
7. (Stage 1: Estimate p -moment to determine number of samples for stage 2.) Create k many blocks. For each block:
 1. Get m samples from the stream.
 2. Add the samples and divide by m to get this block's sample mean, *mean*.
 3. Calculate the p -moment estimate for this block, which is: $(\sum_{i=0, \dots, k-1} (\text{block}[i] - \text{mean})^p)/m$, where $\text{block}[i]$ is the sample at position i of the block (positions start at 0).
8. (Find the median of the p -moment estimates.) Sort the p -moment estimates from step 7 in ascending order, and set *median* to the value in the middle of the sorted list (at position $\text{floor}(k/2)$ with positions starting at 0); this works because k is odd.
9. (Calculate sample size for the next stage.) Set mp to $\max(1, \text{ceil}(\eta * \text{median}^s))$.

10. (Stage 2: Estimate of the sample mean.) Create kp many blocks. For each block:
 1. Get mp samples from the stream.
 2. Add the samples and divide by mp to get this block's sample mean.
11. (Find the median of the sample means. This is definitely an unbiased estimate of the mean when kp is 1 or 2, but unfortunately, it isn't one for any $kp > 2$.) Sort the sample means from step 10 in ascending order, and return the value in the middle of the sorted list (at position $\text{floor}(kp/2)$ with positions starting at 0); this works because kp is odd.

Note: If the stream of random numbers meets the condition for this algorithm for a given q , p , and κ , then it still meets that condition when those numbers are multiplied by a constant or a constant is added to them.

6 Randomized Integration

Monte Carlo integration is a randomized way to estimate the integral of a function. The adaptive algorithm in this article can be used to estimate an integral of a function $f(Z)$, where Z is an n -dimensional vector chosen at random in the sampling domain. The estimate will come within ε of the true integral with probability $1 - \delta$ or greater, as long as the following conditions are met:

- The q^{th} c.a.m. for $f(Z)$ is finite. That is, $\mathbf{E}[\text{abs}(f(Z) - \mathbf{E}[f(Z)])^q]$ is finite.
- The q^{th} c.a.m.'s q^{th} root is no more than κ times the p^{th} c.a.m.'s p^{th} root, where κ is 1 or greater.

Unfortunately, these conditions may be hard to verify in practice, especially when $f(Z)$ is not known. (In fact, $\mathbf{E}[f(Z)]$, as seen above, is the unknown integral that we seek to estimate.)

For this purpose, each number in the stream of random numbers is generated as follows (see also Kunsch et al.):

1. Set Z to an n -dimensional vector (list of n numbers) chosen at random in the sampling domain, independently of any other choice. Usually, Z is chosen *uniformly* at random this way.
2. Calculate $f(Z)$, and set the next number in the stream to that value.

The following example (coded in Python for the SymPy computer algebra library) shows how to find parameter κ for estimating the integral of $\min(Z1, Z2)$ where $Z1$ and $Z2$ are each uniformly chosen at random in the interval $[0, 1]$. It assumes $p = 2$ and $q = 4$. (This is a trivial example because we can calculate the integral directly — $1/3$ — but it shows how to proceed for more complicated cases.)

<h1>Distribution of Z1 and Z2</h1>

```
u1=Uniform('U1',0,1)
u2=Uniform('U2',0,1)
<h1>Function to estimate</h1>
```

```
func = Min(u1,u2)
emean=E(func)
p = S(2) # Degree of p-moment
q = S(4) # Degree of q-moment
<h1>Calculate value for kappa</h1>
```

```
kappa = E(Abs(func-emean)**q)**(1/q) / E(Abs(func-emean)**p)**(1/p)
```

```
pprint(Max(1,kappa))
```

7 Requests and Open Questions

Assume we have an *oracle*, or "black box", that produces independent random numbers X with a known lower and upper bound, but an unknown mean. Given this oracle and possibly a second source of randomness (such as unbiased random bits), are there results that give an algorithm that approximates $f(\mathbf{E}[X])$ with a user-specified error bound, ideally using as few samples from the oracle as possible? Here, f is a known function belonging to a given class of functions.

The algorithm should—

- ensure the expected (absolute) error and/or mean squared error is within a user-specified error tolerance, or if that is not possible,
- return an estimate that is within a user-specified tolerance on the absolute error or relative error with probability greater than $1 - \delta$, where δ is user-specified.

The classes of functions f that are of interest are:

- C1: Continuous functions that map $[0, 1]$ to $[0, 1]$ and are polynomially bounded (meaning that $f(x)$ and $1 - f(x)$ are both bounded from below by $\min(x^n, (1-x)^n)$ for some integer n) (Keane and O'Brien 1994)⁽⁶⁾.
- C2: Continuous functions that map $[0, 1]$ to $[0, 1]$ and are not necessarily polynomially bounded.
- C3: Piecewise continuous functions that map $[0, 1]$ to $[0, 1]$.

Note the following:

- The value the algorithm should approximate is $f(\mathbf{E}[X])$, not $\mathbf{E}[f(X)]$, which is different in general. But algorithms for $\mathbf{E}[f(X)]$ are welcome in case it's more difficult.
- The Gamma Bernoulli Approximation scheme as well as the Kunsch algorithm both estimate the mean of X ($\mathbf{E}[X]$), rather than a function of that mean ($f(\mathbf{E}[X])$), as this request asks.
- Some algorithms produce unbiased estimates of $f(\mathbf{E}[X])$, but not with a user-specified error, as this request asks (Jacob and Thiery 2015)⁽⁷⁾.
- Algorithms like the one being asked for here are especially useful because they can help build so-called "[approximate Bernoulli factories](#)", or algorithms that approximately sample the probability $f(\lambda)$ given a coin with probability of heads of λ . In this case, X should follow the Bernoulli distribution.
- It is suspected that the algorithm's performance will depend on the "smoothness" of $f(X)$ (see also (Holtz et al. 2011)⁽⁸⁾).

Also, are there results that estimate $f(\mathbf{E}[X])$ with a user-specified error bound even if X follows a certain kind of unbounded distribution (such as a geometric, exponential, or Poisson distribution, or another distribution with exponential tails)?

8 Notes

- ⁽¹⁾ Huber, M., 2017. A Bernoulli mean estimate with known relative error distribution. *Random Structures & Algorithms*, 50(2), pp.173-182. (preprint in arXiv:1309.5413v2 [math.ST], 2015).
- ⁽²⁾ Feng, J. et al. "Monte Carlo with User-Specified Relative Error." (2016).
- ⁽³⁾ Huber, Mark, and Bo Jones. "Faster estimates of the mean of bounded random variables."

Mathematics and Computers in Simulation 161 (2019): 93-101.

- (4) Huber, Mark, "[An optimal\(\$\epsilon\$, \$\delta\$ \)-approximation scheme for the mean of random variables with bounded relative variance](#)", arXiv:1706.01478, 2017.
- (5) Kunsch, Robert J., Erich Novak, and Daniel Rudolf. "Solvable integration problems and optimal sample size selection." Journal of Complexity 53 (2019): 40-67. Also in <https://arxiv.org/pdf/1805.08637.pdf>.
- (6) Keane, M. S., and O'Brien, G. L., "A Bernoulli factory", ACM Transactions on Modeling and Computer Simulation 4(2), 1994.
- (7) Pierre E. Jacob. Alexandre H. Thiery. "On nonnegative unbiased estimators." Ann. Statist. 43 (2) 769 - 784, April 2015. <https://doi.org/10.1214/15-AOS1311>
- (8) Holtz, O., Nazarov, F., Peres, Y., "New Coins from Old, Smoothly", Constructive Approximation 33 (2011).

9 License

Any copyright to this page is released to the Public Domain. In case this is not possible, this page is also licensed under [Creative Commons Zero](#).